

Spring Boot and Microservice Interview Questions

1) What is the main class in a Spring Boot application?

→ The main class in a Spring Boot application is the entry point and is annotated with `@SpringBootApplication`

```
@SpringBootApplication  
public class MyApplication {  
    public static void main(String[] args) {  
        SpringApplication.run(MyApplication.class, args);  
    }  
}
```

2) What does the `@SpringBootApplication` annotation do?

→ `@SpringBootApplication` is a convenience annotation that combines three annotations

`@Configuration` : (marks the class as a source of bean definitions)

`@EnableAutoConfiguration` : (enables Spring Boot's auto-configuration mechanism)

`@ComponentScan` : (scans the package of the annotated class for Spring components)

3) How do you handle exceptions in Spring Boot?

- You can handle exceptions in Spring Boot using `@ControllerAdvice` and `@ExceptionHandler` annotations to create a global exception handler.

```
@ControllerAdvice
public class GlobalExceptionHandler {
    @ExceptionHandler(ResourceNotFoundException.class)
    public ResponseEntity<ErrorResponse> handleResourceNotFoundException(ResourceNotFoundException ex) {
        ErrorResponse errorResponse = new ErrorResponse("NOT_FOUND", ex.getMessage());
        return new ResponseEntity<>(errorResponse, HttpStatus.NOT_FOUND);
    }
}
```

4) What is Spring Boot Actuator and what are its benefits?

- Spring Boot Actuator provides production-ready features such as **health checks**, **metrics**, and **monitoring** for Spring Boot application.

```
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-actuator</artifactId>
</dependency>
```

```
management.endpoints.web.exposure.include=health,info
```

5) What are Spring Profiles and how do you use them?

- Spring Profiles allow you to segregate parts of your application configuration and make it only available in certain environments.
- You can activate profiles using the **spring.profiles.active** property.

```
# application-dev.properties

spring.datasource.url=jdbc:mysql://localhost:3306/devdb

# application-prod.properties

spring.datasource.url=jdbc:mysql://localhost:3306/proddb
```

6) How do you optimize the startup time of a Spring Boot application in a production environment?

Spring Boot provides various options for optimizing startup time. Key strategies include:

- 1) **Lazy Initialization** : Use `spring.main.lazy-initialization=true` to delay bean initialization until needed.
- 2) **Profile-Specific Configuration** : Separate configurations per environment to avoid unnecessary loading.
- 3) **Component Scanning** : Restrict the scope of component scanning using `@ComponentScan` to only include essential packages.
- 4) **Reduce Bean Creation** : Avoid creating unnecessary beans during startup, especially for time-intensive services.

By reducing the number of beans initialized upfront, you can significantly speed up application startup.

7) Explain the concept of Spring Boot's @ConfigurationProperties with complex objects. How would you handle nested configurations?

- The **@ConfigurationProperties** annotation is a powerful way to map external configuration properties into Java objects
- For complex, nested configurations, Spring Boot can handle hierarchical properties through nested classes.

```
@ConfigurationProperties(prefix = "app")
public class AppConfig {
    private Database database;
    private List<Service> services;

    public static class Database {
        private String url;
        private String username;
        private String password;
    }

    public static class Service {
        private String name;
        private int timeout;
    }
}
```

8) What are the main challenges with distributed tracing in Spring Boot microservices, and how do you implement it?

- Distributed tracing allows tracking requests across multiple microservices.
- The challenges include latency, proper correlation of requests, and aggregating trace data across services.

Solution:

- **Spring Cloud Sleuth:** Automatically instruments Spring Boot applications for distributed tracing.
- **Integration with Zipkin or Jaeger:** Use Sleuth with tools like Zipkin for trace visualization and monitoring.
- **Correlation:** Propagate TracId and SpanId headers for cross-service correlation, ensuring traceability.

9) What is Spring Boot's @Retryable annotation, and how do you fine-tune it for microservices reliability?

- The @Retryable annotation in Spring Boot allows retrying a method call in case of failure.
- This is essential for improving the reliability of services that might experience transient failures (e.g., network timeouts or database issues).

```
@Retryable(value = {IOException.class}, maxAttempts = 5, backoff = @Backoff(delay =
public String fetchData() {
    // API call that might fail
}
```

10) How do you handle versioning in Spring Boot APIs?

API versioning is crucial for maintaining backward compatibility as your services evolve. Common strategies include:

URI Versioning	:	/api/v1/resource
Parameter Versioning	:	/api/resource?version=1
Header-based Versioning	:	Through custom headers like API-Version.

11) How do you manage external configurations in a Spring Boot application across multiple environments?

- **Profiles:** Use `@Profile` to define beans for specific environments (e.g.,
`@Profile("prod")`).
- **application.properties or YAML:** Use `application-prod.properties` for production-specific configurations.
- **Spring Cloud Config:** For distributed systems, use Spring Cloud Config Server to manage configurations centrally.

12) What are some common performance bottlenecks in Spring Boot applications and how do you resolve them?

1. Database Access

- **Optimize Queries:** Ensure SQL queries are efficient and use proper indexes.
 - **Pagination:** Fetch data in chunks using pagination to avoid memory overload.
 - **Connection Pooling:** Leverage tools like **HikariCP** for efficient database connection management.
-

2. Memory Management

- **Avoid Memory Leaks:** Regularly review your code for unclosed resources or lingering references.
 - **Monitoring:** Use tools like **VisualVM** or **JVisualVM** to profile memory usage and identify leaks.
-

3. Thread Pooling

- **Thread Pool Sizing:** Tune thread pools based on system resources and expected load.
 - **Executor Services:** Use Java's ExecutorService or Spring's @Async with properly configured task executors for managing background tasks and HTTP request handling.
 -
-

4. Caching

- **Spring Caching Abstraction:** Use annotations like @Cacheable, @CachePut, and @CacheEvict to reduce database hits.
- **Cache Providers:** Integrate with **Ehcache**, **Caffeine**, or **Redis** for high-performance caching solutions.

13) How do you configure and manage Spring Boot logging in production?

1. Logback (Default Logging Framework)

-  **Out-of-the-Box Support:** Spring Boot uses **Logback** as the default logging implementation.
-  **Custom Configuration:** Configure advanced logging using `logback-spring.xml` for environment-specific setups.

2. External Logging Integration

-  **Centralized Logging:** Integrate with logging stacks like ELK (Elasticsearch + Logstash + Kibana) or EFK (Fluentd instead of Logstash).
-  **Benefits:** Enables real-time log aggregation, analysis, and visualization across distributed services.

3. Log Levels per Environment

-  **Environment-Specific Tuning:**
 - DEBUG – for development to get detailed application flow.
 - INFO – for production to reduce noise but keep essential logs.
-  **Configuration:** Set log levels in `application.yml` or `application.properties`:

14) How can you disable a specific auto-configuration class in Spring Boot ?

→ Disable an auto-configuration class by adding the exclude attribute:

```
@SpringBootApplication(exclude = { DataSourceAutoConfiguration.class })
public class MyApplication {
    public static void main(String[] args) {
        SpringApplication.run(MyApplication.class, args);
    }
}
```

```
@EnableAutoConfiguration(exclude = { DataSourceAutoConfiguration.class })
public class MyApplication {
    public static void main(String[] args) {
        SpringApplication.run(MyApplication.class, args);
    }
}
```

15) How to configure multiple datasource in spring boot application ?

→ Define Properties in application.yml or application.properties .

```
spring:  
  datasource:  
    primary:  
      url: jdbc:mysql://localhost:3306/primary_db  
      username: root  
      password: password  
      driver-class-name: com.mysql.cj.jdbc.Driver  
  
    secondary:  
      url: jdbc:h2:mem:secondary_db  
      driver-class-name: org.h2.Driver  
      username: sa  
      password:
```

➔ Configure primary datasource

```
@Configuration
@EnableTransactionManagement
@EnableJpaRepositories(
    basePackages = "com.example.repo.primary",
    entityManagerFactoryRef = "primaryEntityManagerFactory",
    transactionManagerRef = "primaryTransactionManager"
)
public class PrimaryDataSourceConfig {

    @Bean
    @Primary
    @ConfigurationProperties("spring.datasource.primary")
    public DataSource primaryDataSource() {
        return DataSourceBuilder.create().build();
    }

    @Primary
    @Bean
    public LocalContainerEntityManagerFactoryBean primaryEntityManagerFactory(
        EntityManagerFactoryBuilder builder) {
        return builder
            .dataSource(primaryDataSource())
            .packages("com.example.entity.primary")
            .persistenceUnit("primary")
            .build();
    }

    @Primary
    @Bean
    public PlatformTransactionManager primaryTransactionManager(
        EntityManagerFactory primaryEntityManagerFactory) {
        return new JpaTransactionManager(primaryEntityManagerFactory);
    }
}
```

Configure secondary datasource

```
@Configuration
@EnableJpaRepositories(
    basePackages = "com.example.repo.secondary",
    entityManagerFactoryRef = "secondaryEntityManagerFactory",
    transactionManagerRef = "secondaryTransactionManager"
)
public class SecondaryDataSourceConfig {

    @Bean
    @ConfigurationProperties("spring.datasource.secondary")
    public DataSource secondaryDataSource() {
        return DataSourceBuilder.create().build();
    }

    @Bean
    public LocalContainerEntityManagerFactoryBean secondaryEntityManagerFactory(
        EntityManagerFactoryBuilder builder) {
        return builder
            .dataSource(secondaryDataSource())
            .packages("com.example.entity.secondary")
            .persistenceUnit("secondary")
            .build();
    }

    @Bean
    public PlatformTransactionManager secondaryTransactionManager(
        EntityManagerFactory secondaryEntityManagerFactory) {
        return new JpaTransactionManager(secondaryEntityManagerFactory);
    }
}
```

Key Points

- Use `@Primary` to mark one datasource as the default.
- Define separate `basePackages` for `repositories` and `entities` for each data source.
- Each data source must have its own:
 - `DataSource`
 - `EntityManagerFactory`
 - `TransactionManager`

16) Can you explain the purpose of Stereotype annotations in the Spring Framework ?

- Stereotype annotations in Spring are specialized annotations used to declare Spring-managed components.
- They help Spring automatically detect and register beans during component scanning.

💡 Why Use Stereotype Annotations?

1. Enable Component Scanning
 - Classes annotated with stereotypes are automatically discovered and registered as beans.
2. Improve Code Semantics
 - Using `@Service`, `@Repository`, and `@Controller` makes your code more readable and meaningful by indicating the role of the class.
3. Enable Aspect-Oriented Features
 - For example, `@Repository` enables exception translation using Spring's `PersistenceExceptionTranslationPostProcessor`.
4. Reduce Boilerplate Configuration
 - No need for explicit `<bean>` declarations in XML or manual `@Bean` registration in config classes.

17) How many ways we can perform dependency injection in spring or spring boot ?

1. Constructor-Based Injection (Recommended ✓)

- Dependencies are injected via the class constructor.

```
java  
  
@Service  
public class OrderService {  
    private final PaymentService paymentService;  
  
    @Autowired  
    public OrderService(PaymentService paymentService) {  
        this.paymentService = paymentService;  
    }  
  
    // Or even better (since Spring 4.3+), no need for @Autowired  
    public OrderService(PaymentService paymentService) {  
        this.paymentService = paymentService;  
    }  
}
```

Benefits:

- Promotes immutability
- Easy for unit testing
- Preferred for mandatory dependencies

2. Setter-Based Injection

- Dependencies are injected via setter methods.

```
java  
  
@Service  
public class OrderService {  
  
    private PaymentService paymentService;  
  
    @Autowired  
    public void setPaymentService(PaymentService paymentService) {  
        this.paymentService = paymentService;  
    }  
}
```

Use Case:

- When the dependency is optional or needs to be set after construction

3. Field-Based Injection (Not recommended ✗)

- Dependencies are injected directly into fields.

```
java  
  
@Service  
public class OrderService {  
  
    @Autowired  
    private PaymentService paymentService;  
}
```

Drawbacks:

- Harder to test (no constructor for mocking)
- Violates principles of encapsulation
- Not ideal for required dependencies

18) Can you provide an example of a real-world use case where `@PostConstruct` is particularly useful?

→ Preloading Data on Application Startup

```
Copy Edit  
@Service  
public class ProductCacheService {  
  
    private final ProductRepository productRepository;  
    private Map<Long, Product> productCache = new HashMap<>();  
  
    public ProductCacheService(ProductRepository productRepository) {  
        this.productRepository = productRepository;  
    }  
  
    @PostConstruct  
    public void initCache() {  
        List<Product> products = productRepository.findAll();  
        products.forEach(product -> productCache.put(product.getId(), product));  
        System.out.println("✓ Product cache initialized with " + products.size() + " items");  
    }  
  
    public Product getProductById(Long id) {  
        return productCache.get(id);  
    }  
}
```

💡 What Happens Here?

- Spring injects the `ProductRepository`.
- After the bean is fully constructed and dependencies injected, `@PostConstruct` runs.
- The `initCache()` method loads all product data into a local cache for faster access.

💡 Why Use `@PostConstruct` Instead of a Constructor?

- You can't call other Spring beans safely in a constructor.
- `@PostConstruct` ensures the bean is fully initialized before the method runs.

✓ Summary

Use Case	Why <code>@PostConstruct</code> Helps
Load data into cache	Waits until all dependencies are ready
Register with external services	Avoids initialization race conditions
Validate bean configuration	Ensures consistent startup checks
Trigger scheduled jobs manually	Pre-load cron logic if needed

19) How can we dynamically load values in a Spring Boot application?

1. Using @Value Annotation (for simple values)

```
@Component
public class AppInfo {

    @Value("${app.name}")
    private String appName;

    public void printAppName() {
        System.out.println("App Name: " + appName);
    }
}
```

2. Using @ConfigurationProperties (for grouped values)

```
yaml
app:
  config:
    name: Inventory System
    version: 1.0
    enabled: true

java
@Component
@ConfigurationProperties(prefix = "app.config")
public class AppConfig {
    private String name;
    private String version;
    private boolean enabled;

    // Getters and setters
}
```

3. Environment Object (for runtime value resolution)

```
@Component
public class DynamicLoader {

    @Autowired
    private Environment env;

    public void load() {
        String dbUser = env.getProperty("spring.datasource.username");
        System.out.println("Database User: " + dbUser);
    }
}
```

Good for runtime evaluation or conditional config logic.

4. External Configuration (via Environment Variables or CLI)

```
java -jar app.jar --server.port=8085 --app.name="Staging App"
```

5. Dynamic Loading from a Database / Remote Source

→ You can fetch values at runtime (e.g., for feature flags or settings) using `@PostConstruct`, a scheduled task, or an external config service.

```
@Service
public class ConfigService {

    private final ConfigRepository configRepository;

    private Map<String, String> dynamicConfig = new HashMap<>();

    public ConfigService(ConfigRepository configRepository) {
        this.configRepository = configRepository;
    }

    @PostConstruct
    public void loadFromDb() {
        List<ConfigEntity> configs = configRepository.findAll();
        configs.forEach(c -> dynamicConfig.put(c.getKey(), c.getValue()));
    }

    public String getValue(String key) {
        return dynamicConfig.getOrDefault(key, "N/A");
    }
}
```

20) How to load External Properties in Spring Boot ?

→ Loading external properties in a Spring Boot application is a common way to separate configuration from code — especially useful for different environments (dev, test, prod).

1. Use --spring.config.location Command-Line Argument

```
java -jar myapp.jar --spring.config.location=file:/path/to/config.properties
```

2. Use SPRING_CONFIG_LOCATION Environment Variable

```
export SPRING_CONFIG_LOCATION=file:/path/to/config/
java -jar myapp.jar
```

● Useful in cloud or containerized environments (e.g., Kubernetes, Docker).

3. Use @PropertySource in Java Config (for custom file names)

```
@Configuration
@PropertySource("file:/external/config/custom-config.properties")
public class ExternalConfig { }
```

- Works with .properties, but not .yml
- Use Environment or @Value to access values

4. Use External Config in Spring Boot Profiles

Structure external configs per environment:

```
arduino                                     ⌂ Copy  
/config/  
    ├── application.yml  
    ├── application-dev.yml  
    └── application-prod.yml
```

Then activate the profile:

```
bash                                         ⌂ Copy  
java -jar app.jar --spring.profiles.active=dev --spring.config.location=file:/config/
```

21) Can we avoid this dependency ambiguity without using @Qualifier ?

- ◆ 1. Use @Primary to Designate a Default Bean

```
@Bean  
@Primary  
public NotificationService emailNotificationService() {  
    return new EmailNotificationService();  
}  
  
@Bean  
public NotificationService smsNotificationService() {  
    return new SmsNotificationService();  
}
```

- ◆ 2. Inject by Concrete Class Instead of Interface

```
@Autowired  
private SmsNotificationService smsNotificationService;
```

- ◆ 3. Use **@ComponentScan** Selectively

Avoid registering multiple beans of the same type by refining your **@ComponentScan**.

```
@SpringBootApplication  
@ComponentScan(basePackages = "com.example.email") // excludes other implementations  
public class MyApp {}
```

- ◆ 4. Use Profiles to Load Only One Bean at a Time

```
java  
  
@Bean  
@Profile("email")  
public NotificationService emailNotificationService() {  
    return new EmailNotificationService();  
}  
  
@Bean  
@Profile("sms")  
public NotificationService smsNotificationService() {  
    return new SmsNotificationService();  
}
```

```
--spring.profiles.active=email
```

22) What is bean scope & Can you explain different type of bean scope ?

Common Types of Bean Scopes in Spring

Spring supports several scopes, especially in different environments (web vs. non-web).

Scope	Description	Applicable In
singleton	Only one instance per Spring container (default)	Core, Web
prototype	New instance each time it is requested	Core, Web
request	One instance per HTTP request	Web (Servlet)
session	One instance per HTTP session	Web (Servlet)
application	One instance per ServletContext (application-wide)	Web (Servlet)
websocket	One instance per WebSocket session	WebSocket apps

23) Can you provide a few real-time use cases for when to choose Singleton scope and Prototype scope ?

Summary Comparison

Aspect	Singleton Scope	Prototype Scope
Instance Count	One per Spring container	New every time it's requested
Lifecycle Managed	Fully by Spring	Only initialization (no destroy callback)
Use For	Stateless, shared, thread-safe beans	Stateful, non-thread-safe, custom state
Examples	Service, Repository, Config, Caching	Report builder, Upload job, Worker task

24) How can we deserialize a JSON request payload into an object ?

→ Jackson

25) How can we perform content negotiation (XML/JSON) in Rest endpoint ?

1. Using HTTP Headers (Most common way)

Client sends:

- `Accept` header indicating the desired response format.
 - Example:
 - `Accept: application/json`
 - `Accept: application/xml`
 - Or `Accept: application/json, application/xml;q=0.8` (priority-based)

Server responds:

- Examines the `Accept` header.
- Returns the response in the requested format.
- Sets `Content-Type` header accordingly (`application/json` or `application/xml`).

2. How to implement in popular frameworks

In Spring Boot (Java)

- Just add support for both JSON and XML in your `pom.xml` or `build.gradle` (Jackson for JSON and JAXB or Jackson XML for XML).
- Use `@RestController` and simply return the POJO.
- Spring Boot automatically negotiates content based on the `Accept` header.

26) How do you maintain the versioning for your REST API?

1. URI Path Versioning (Most common & explicit)

Add the version in the URL path:

```
bash  
  
/api/v1/users  
/api/v2/users
```

2. Request Header Versioning

Use a custom header (e.g., `Accept` or `X-API-Version`):

```
bash  
  
GET /api/users  
Accept: application/vnd.myapp.v1+json
```

3. Query Parameter Versioning

Pass version as a query param:

```
pgsql  
  
GET /api/users?version=1
```

27) How will you document your rest API ?

Ways to Document REST APIs

1. Use OpenAPI (Swagger) Specification

- The industry standard for REST API documentation.
- Describes endpoints, request/response formats, status codes, authentication, etc.
- Can generate interactive documentation (Swagger UI) where users can try out API calls.

3. Write Manual Documentation

- Create human-readable docs that explain:
 - Available endpoints and their HTTP methods
 - Required/request parameters and body
 - Response formats and status codes
 - Error messages
 - Authentication/authorization
- Can be done in markdown, wikis, or API portals like [ReadMe](#), [Postman](#), or [GitHub Pages](#).

2. Auto-generate Docs From Code Annotations

- Use tools that scan your code and produce OpenAPI specs automatically.
- Examples:
 - **Spring Boot:** Use `springdoc-openapi` or `springfox-swagger2`.
 - **Express:** Use `swagger-jsdoc` and `swagger-ui-express`.
 - **Flask:** Use `flasgger` or `apispec`.

28) How will you consume restful API ?

👤 Java (Spring Boot with `RestTemplate` or `WebClient`)

Using `RestTemplate` (Legacy)

```
java

RestTemplate restTemplate = new RestTemplate();
String url = "https://api.example.com/users/1";
User user = restTemplate.getForObject(url, User.class);
```

Using `WebClient` (Preferred for reactive & modern apps)

```
java

WebClient webClient = WebClient.create();

User user = webClient.get()
    .uri("https://api.example.com/users/1")
    .retrieve()
    .bodyToMono(User.class)
    .block(); // blocking call for synchronous processing
```

29) How can you hide certain REST endpoints to prevent them from being exposed externally?

2. Do Not Document Them in OpenAPI/Swagger

Exclude endpoints from appearing in your Swagger UI or OpenAPI spec.

Spring Boot Example with Springdoc OpenAPI:

```
java

@Hidden
@GetMapping("/api/internal/debug")
public String debugEndpoint() {
    return "Internal info";
}
```

1. Use Spring Security (or any auth framework)

Restrict access to internal endpoints by role or IP.

Spring Boot Example:

```
java

@Override
protected void configure(HttpSecurity http) throws Exception {
    http
        .authorizeRequests()
        .antMatchers("/api/internal/**").hasRole("ADMIN") // Only ADMIN can access
        .antMatchers("/api/public/**").permitAll(); // Publicly accessible
}
```

5. Use Profiles or Feature Flags

Only expose certain controllers or endpoints in specific environments (e.g., dev).

Spring Boot Example:

```
java

@Profile("dev")
@RestController
@RequestMapping("/api/internal")
public class InternalController {
    // only available in dev profile
}
```

30) How can you avoid defining handlers for multiple exceptions, or what is the best practice for handling exceptions ?

→ Instead of writing try-catch blocks or separate handlers in every controller , use a **centralized exception handler** to handle all exceptions in one place.

```
@RestControllerAdvice
public class GlobalExceptionHandler {

    @ExceptionHandler(ResourceNotFoundException.class)
    public ResponseEntity<ErrorResponse> handleResourceNotFound(ResourceNotFoundException ex) {
        ErrorResponse error = new ErrorResponse(
            HttpStatus.NOT_FOUND.value(),
            ex.getMessage(),
            LocalDateTime.now()
        );
        return new ResponseEntity<>(error, HttpStatus.NOT_FOUND);
    }

    @ExceptionHandler(Exception.class)
    public ResponseEntity<ErrorResponse> handleGenericException(Exception ex) {
        ErrorResponse error = new ErrorResponse(
            HttpStatus.INTERNAL_SERVER_ERROR.value(),
            "An unexpected error occurred",
            LocalDateTime.now()
        );
        return new ResponseEntity<>(error, HttpStatus.INTERNAL_SERVER_ERROR);
    }
}
```

```
public class ErrorResponse {
    private int status;
    private String message;
    private LocalDateTime timestamp;

    // constructors, getters, setters
}
```

✓ Benefits of Centralized Exception Handling:

- No repetitive try-catch blocks.
- Cleaner controller code.
- Consistent error response format.
- Easier to add logging, monitoring, alerts.

- Group similar exceptions if needed:

```
java Copy
@ExceptionHandler({IllegalArgumentException.class, IllegalStateException.class})
public ResponseEntity<?> handleBadRequest(RuntimeException ex) { ... }
```

31) How will you validate or sanitise your input payload ?

1. Use Bean Validation

Use annotations like @NotNull, @Email, @Size, etc., in your DTOs with javax.validation or jakarta.validation.

```
public class UserRequest {  
  
    @NotBlank(message = "Name is required")  
    private String name;  
  
    @Email(message = "Email must be valid")  
    private String email;  
  
    @Min(value = 18, message = "Age must be at least 18")  
    private int age;  
  
    // getters and setters  
}
```

In Controller:

```
java  
  
@PostMapping("/users")  
public ResponseEntity<?> createUser(@Valid @RequestBody UserRequest request)  
    // Validated input  
    return ResponseEntity.ok("User created");  
}
```

2. Handle Validation Errors Gracefully

Use `@RestControllerAdvice` to return consistent messages.

```
java  
  
@ExceptionHandler(MethodArgumentNotValidException.class)  
public ResponseEntity<?> handleValidationErrors(MethodArgumentNotValidException ex){  
    Map<String, String> errors = new HashMap<>();  
    ex.getBindingResult().getFieldErrors().forEach(error ->  
        errors.put(error.getField(), error.getDefaultMessage()));  
    return ResponseEntity.badRequest().body(errors);  
}
```

3. Sanitize Input (Optional, Based on Risk)

Sanitization is usually about cleaning or escaping input before using it (e.g., in UI or SQL).

- Use frameworks (Hibernate, Spring Data JPA) to prevent SQL injection.
- For strings, you can strip/escape harmful HTML or script content.

Example: Preventing XSS (if rendering user input in UI)

```
java  
  
String cleanInput = HtmlUtils.htmlEscape(userInput);
```

32) How can you populate validation error message to the end users ?

→ To populate validation error messages to end users, especially in REST APIs (like with Spring Boot), you need to:

- 1) Use validation annotations in your DTO.
- 2) Catch validation errors centrally.
- 3) Return a structured, user-friendly error response.

◆ 1. Add Validation Annotations in DTO

```
public class UserRequest {  
  
    @NotBlank(message = "Name is required")  
    private String name;  
  
    @Email(message = "Please provide a valid email")  
    private String email;  
  
    @Min(value = 18, message = "Age must be at least 18")  
    private int age;  
  
    // getters and setters  
}
```

2. Use @Valid in controller

```
@PostMapping("/users")  
public ResponseEntity<?> createUser(@Valid @RequestBody UserRequest request) {  
    // Only called if validation passes  
    return ResponseEntity.ok("User created");  
}
```

3 . Handle validation with `@ControllerAdvice`

```
@RestControllerAdvice  
public class GlobalExceptionHandler {  
  
    @ExceptionHandler(MethodArgumentNotValidException.class)  
    public ResponseEntity<Map<String, String>> handleValidationErrors(MethodArgumentNotValidException ex) {  
        Map<String, String> errors = new HashMap<>();  
        ex.getBindingResult().getFieldErrors().forEach(error ->  
            errors.put(error.getField(), error.getDefaultMessage())  
        );  
        return new ResponseEntity<>(errors, HttpStatus.BAD_REQUEST);  
    }  
}
```

💡 Best Practices

Practice	Benefit
Use custom error messages	More user-friendly
Centralize validation handling	DRY code and consistent error format
Return clear field-error mapping	Helps frontend show exact field issues
Localize messages (optional)	Support for international users

🌐 Bonus: Localization Support

Add to `messages.properties`:

properties

user.email.invalid=Please provide a valid email
user.name.blank=Name is required

Then update annotations:

java

```
@Email(message = "{user.email.invalid}")  
@NotBlank(message = "{user.name.blank}")
```

33) Let's say you find a bug in production environment and now you want to debug that scenario ,How can you do that from your local ?

→ Debugging a bug found in the production environment from your local can be tricky.

◆ 1. Understand and Reproduce the Bug

- Gather details from logs, bug reports, monitoring tools, or support tickets:
 - Error messages
 - Request payloads
 - Stack traces
 - Affected users
 - Timestamp of issue
- Ask: Can the issue be reproduced reliably?

◆ 2. Check Production Logs / Monitoring Tools

Use observability tools such as:

- 🔎 ELK Stack (Elasticsearch, Logstash, Kibana)
- 📈 Prometheus + Grafana
- 📁 Log files via CloudWatch (AWS), Stackdriver (GCP), or other logging systems.

Get the exact:

- Input/request data
- Headers
- Auth tokens (if applicable)
- DB state (snapshot or relevant rows)

◆ 3. Mimic Production Environment Locally

- Checkout the same Git branch or tag deployed to prod.
- Configure the same:
 - Environment variables
 - App configs (like external services, feature flags)
- Use Docker / Docker Compose / Kubernetes manifests to simulate the same infrastructure locally.

◆ 4. Replay the Bug Locally

- Use the exact same input from production (e.g., REST API request).
- Tools:
 - Postman (to replicate requests)
 - `curl` with prod payloads
 - JUnit/Integration tests with prod data
- Simulate edge cases like:
 - Null/missing fields
 - Invalid states in DB

◆ 5. Attach Debugger

In your IDE:

- Set breakpoints around the suspect code.
- Run the app in **debug mode** (`debug` profile or IDE debug button).
- Step through the flow using actual data that caused the issue.

◆ 7. Fix the Bug & Retest Locally

- Once identified, fix the issue.
- Rerun the test.
- Run full test suite (unit + integration + contract tests).

◆ 8. Peer Review + Deploy Safely

- Create a **pull request** and request **code review**.
- Deploy to staging or QA first, verify the fix.
- Use **feature toggles** or **canary releases** to roll out gradually if needed.

◆ 6. Write Failing Test Case (TDD-style Debugging)

Write a unit or integration test using the same input and DB state. This helps:

- Isolate the problem.
- Prevent future regressions (part of automated test suite).

java

```
@Test  
void shouldThrowException_WhenInputIsInvalid() {  
    // simulate production scenario  
}
```

 Copy  Edit

◆ 9. Monitor After Deployment

- Watch logs and alerts post-deployment.
- Confirm the issue no longer occurs.
- Possibly add observability (metrics, logs) in fixed area.

34) What is the key benefit of using the Spring Data JPA ?

1. 🚀 Reduces Boilerplate Code

- No need to write repetitive CRUD logic.
- You get basic methods like `save()`, `findById()`, `findAll()`, `delete()` out of the box

```
java
```

```
public interface UserRepository extends JpaRepository<User, Long> {}
```

3. 📋 Supports JPQL and Native Queries

- When custom logic is needed, you can still write JPQL or native SQL.

```
java
```

```
@Query("SELECT u FROM User u WHERE u.age > :age")
List<User> findUsersOlderThan(@Param("age") int age);
```

9. 🔍 Repository Customization

- You can define your own methods alongside Spring-generated ones.
- Supports custom repository implementation for complex logic.

2. 💬 Powerful Query Generation from Method Names

- You can define queries just by naming methods in a certain way.

```
java
```

```
List<User> findByEmail(String email);
List<User> findByFirstNameAndLastName(String first, String last);
```

Spring automatically translates these into SQL queries.

4. 📦 Pagination and Sorting Support

- Built-in support for paging and sorting results.

```
java
```

```
Page<User> findAll(Pageable pageable);
List<User> findAll(Sort sort);
```

35) What are the different ways to define custom queries in Spring Data JPA ?

◆ 1. Derived Query Methods (Query by Method Name)

Let Spring generate the SQL based on method name.

```
java  
  
List<User> findByFirstNameAndStatus(String firstName, String status);
```

✓ Simple and readable, but limited for complex queries.

◆ 3. Native SQL with `@Query(nativeQuery = true)`

When you need full SQL power (joins, DB-specific functions).

```
java  
  
@Query(value = "SELECT * FROM users WHERE email = :email", nativeQuery = true)  
User findByEmailNative(@Param("email") String email);
```

✓ Useful when JPQL doesn't support required operations.

◆ 2. JPQL with `@Query` Annotation

Use JPQL (object-oriented SQL) to define custom queries.

```
java  
  
@Query("SELECT u FROM User u WHERE u.status = :status")  
List<User> findActiveUsers(@Param("status") String status);
```

✓ Uses entity field names, not table/column names.

◆ 4. Named Queries (Defined in Entity)

Define queries inside entity using `@NamedQuery`.

```
java  
  
@Entity  
 @NamedQuery(name = "User.findByEmail", query = "SELECT u FROM User u WHERE u.email = :email")  
public class User { ... }  
  
interface UserRepository extends JpaRepository<User, Long> {  
    User findByEmail(@Param("email") String email); // Will use the named query  
}
```

✓ Good for reuse across multiple places, but harder to maintain.

◆ 7. Custom Repository Implementation

When logic is too complex for annotations or specifications.

```
java  
  
public interface UserRepositoryCustom {  
    List<User> customSearchLogic(String keyword);  
}  
  
public class UserRepositoryImpl implements UserRepositoryCustom {  
    @PersistenceContext  
    private EntityManager em;  
  
    @Override  
    public List<User> customSearchLogic(String keyword) {  
        return em.createQuery("SELECT u FROM User u WHERE u.name LIKE :kw", User.class)  
            .setParameter("kw", "%" + keyword + "%")  
            .getResultList();  
    }  
}
```

◆ 6. Dynamic Queries with Specification<T> (Criteria API)

Use `JpaSpecificationExecutor<T>` for building dynamic queries at runtime.

```
java  
  
public class UserSpecs {  
    public static Specification<User> hasStatus(String status) {  
        return (root, query, cb) -> cb.equal(root.get("status"), status);  
    }  
}
```

Usage:

```
java  
  
userRepository.findAll(UserSpecs.hasStatus("ACTIVE"));
```

✓ Best for dynamic search filters.

36) Relationship mapping ?

Relationship Type	Annotation	Example Entities	Notes
One-to-One	<code>@OneToOne</code>	User → Profile	Use <code>@JoinColumn</code>
One-to-Many	<code>@OneToMany</code>	Department → Employees	Use <code>mappedBy</code> on parent side
Many-to-One	<code>@ManyToOne</code>	Employee → Department	Child owns the FK
Many-to-Many	<code>@ManyToMany</code>	Student ↔ Course	Uses a join table

◆ 1. Customer Entity (Parent)

java

Copy

```
@Entity
public class Customer {

    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    private String name;
    private String email;

    @OneToMany(mappedBy = "customer", cascade = CascadeType.ALL, orphanRemoval = true)
    private List<Order> orders = new ArrayList<>();

    // Getters and Setters
}
```

◆ 2. Order Entity (Child)

java

```
@Entity
@Table(name = "orders") // optional: avoids SQL reserved keyword conflict
public class Order {

    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    private LocalDate orderDate;
    private Double totalAmount;

    @ManyToOne
    @JoinColumn(name = "customer_id") // foreign key in orders table
    private Customer customer;

    // Getters and Setters
}
```

🧠 Explanation of Key Points

Concept	Applied On	Explanation
@OneToMany	Customer.orders	One customer has many orders
mappedBy = "customer"	Tells JPA that Order.customer owns the relationship	
@ManyToOne	Order.customer	Many orders can belong to one customer
@JoinColumn	On Order.customer	Adds a customer_id column in orders table
CascadeType.ALL	On Customer.orders	Saves/deletes orders when customer is saved/deleted
orphanRemoval = true	On Customer.orders	Deletes orders removed from the list in code

37) Is this possible to execute Join query in Spring Data JPA ? If yes, how can you add some insights ?

◆ 1. JPQL Join (Preferred for Entity Relationships)

Example: Join `Customer` and `Order`

Assuming:

- A `Customer` entity has `@OneToMany` Orders
- An `Order` has `@ManyToOne` Customer

java

```
@Query("SELECT o FROM Order o JOIN o.customer c WHERE c.email = :email")
List<Order> findOrdersByCustomerEmail(@Param("email") String email);
```

This will fetch orders placed by a customer with the specified email.

◆ 2. Join with Fetch (To Avoid N+1 Problem)

Fetch joins load associated entities in a single query.

java

```
@Query("SELECT c FROM Customer c JOIN FETCH c.orders WHERE c.id = :id")
Customer findCustomerWithOrders(@Param("id") Long id);
```

⚡ Improves performance by preventing multiple SQL queries.

◆ 3. Native SQL Join (When You Need Table/Column Control)

You can write raw SQL if needed.

java

```
@Query(value = "SELECT o.* FROM orders o INNER JOIN customer c ON o.customer_id
List<Order> findOrdersByCustomerEmailNative(@Param("email") String email);
```

✓ Best when you need full control over complex joins or use SQL-specific features.

38) How will you implement pagination & Sorting in Spring Data JPA ?

◆ 1. Repository Setup

Your repository must extend `JpaRepository` (or `PagingAndSortingRepository`):

```
java

public interface ProductRepository extends JpaRepository<Product, Long> {
    Page<Product> findByCategory(String category, Pageable pageable);
}
```

`Pageable` enables pagination & sorting in a single argument.

◆ 3. Sorting Alone (without pagination)

```
java

List<Product> products = productRepository.findAll(Sort.by("name").ascending());
```

🔧 Sorting on Multiple Fields

```
java

Sort sort = Sort.by(Sort.Order.desc("price"), Sort.Order.asc("name"));
Pageable pageable = PageRequest.of(0, 10, sort);
```

◆ 2. Using `Pageable` in Service or Controller

► Paginate all products (with sorting)

```
java

Pageable pageable = PageRequest.of(pageNumber, pageSize, Sort.by("price").descending());
Page<Product> pagedResult = productRepository.findAll(pageable);
```

- `PageRequest.of(page, size)` → 0-based page index
- `Sort.by("field").direction()` → ascending or descending

⌚ Custom Queries with Pagination

```
java

@Query("SELECT p FROM Product p WHERE p.category = :category")
Page<Product> findByCategory(@Param("category") String category, Pageable pageable);
```

Works seamlessly with both JPQL and native queries.

39) What is transaction management and how it works ? explain

→ Transaction Management in Spring (or any enterprise app) ensures that a set of operations either all succeed or all fail as one atomic unit of work.

How Transaction Management Works in Spring

→ Spring provides declarative transaction management using the @Transactional annotation.

◆ 1. Using @Transactional Annotation

```
java

@Service
public class OrderService {

    @Autowired
    private OrderRepository orderRepository;

    @Transactional
    public void placeOrder(Order order) {
        orderRepository.save(order);
        // if any exception occurs here, everything rolls back
    }
}
```

All DB operations within this method are part of one transaction. If any exception occurs, everything is rolled back automatically.

◆ 2. Transactional Boundaries

Spring creates a proxy around your bean. When a method marked with `@Transactional` is called from outside the bean, Spring begins a transaction, manages it, and commits/rolls back based on outcome.

◆ 3. Transaction Rollback Rules

By default:

- Unchecked exceptions (`RuntimeException`, `Error`) → trigger rollback
- Checked exceptions → do not trigger rollback unless explicitly configured:

```
java
```

```
@Transactional(rollbackFor = Exception.class)
```

◆ 5. Isolation Levels

Controls visibility of data between transactions.

Level	Description
READ_UNCOMMITTED	Can read uncommitted changes (dirty read)
READ_COMMITTED	Only reads committed data
REPEATABLE_READ	Ensures same data is read again
SERIALIZABLE	Fully isolated (slowest but safest)

Example:

```
java
```

```
@Transactional(isolation = Isolation.READ_COMMITTED)
```

◆ 4. Propagation Behavior

Defines how a transaction behaves when a method is called from another transactional method.

Propagation	Meaning
REQUIRED (default)	Join existing or create a new one
REQUIRES_NEW	Always creates a new transaction
NESTED	Executes within a nested transaction
MANDATORY	Must run inside a transaction, or fail
NOT_SUPPORTED	Runs without a transaction

Real-Life Example

```
java
```

```
@Transactional
public void transferMoney(Long fromAccount, Long toAccount, BigDecimal amount) {
    Account source = accountRepo.findById(fromAccount).get();
    Account target = accountRepo.findById(toAccount).get();

    source.debit(amount);
    target.credit(amount);

    accountRepo.save(source);
    accountRepo.save(target);
}
```

If any step fails, no money is transferred.

40) Can we use transaction in private methods ?

→ No, you cannot use @Transactional effectively on private methods in Spring.

41) How will you handle transaction in distributed microservice ?

- Handling transactions in distributed microservices is one of the most challenging aspects of microservices architecture.
- Since each microservice has its own database, traditional ACID transactions don't work across services.
- Therefore, we use eventual consistency and patterns like

1. Saga Pattern (Recommended)

The Saga Pattern manages distributed transactions using a sequence of local transactions. Each service performs its task and then publishes an event to trigger the next step.

- ◆ Two Types of Sagas:

Type	Description
Choreography	No central coordinator. Services listen and react to events.
Orchestration	A central "orchestrator" tells services what to do next.

3. Event-Driven Architecture (EDA)

Use Kafka, RabbitMQ, or EventBridge to emit events after local DB commits. Other services consume those events and act accordingly.

- Use Outbox Pattern to safely publish events after DB commits.
- Use Change Data Capture (CDC) with Debezium.

42) Explain orchestration and choreography pattern with example ?

🌀 1. Choreography-Based Saga (Event-Driven)

Flow:

Each service listens to events and publishes the next event.

No central control.

Use Case: E-commerce Order Placement

Microservices involved:

- 1) Order Service
- 2) Inventory Service
- 3) Payment Service

◆ **Dependencies:**

- Spring Boot
- Kafka (or RabbitMQ)

◆ Sample Events (DTOs)

```
java

public class OrderCreatedEvent {
    private Long orderId;
    private List<Long> productIds;
}

public class InventoryReservedEvent {
    private Long orderId;
}

public class PaymentCompletedEvent {
    private Long orderId;
}
```

◆ Order Service

```
java

@Service
public class OrderService {

    @Autowired
    private KafkaTemplate<String, Object> kafkaTemplate;

    public void placeOrder(Order order) {
        // save order to DB
        kafkaTemplate.send("order-created", new OrderCreatedEvent(order.getId(), order.getProductIds()));
    }
}
```

◆ Order Service (final completion)

```
java

@KafkaListener(topics = "payment-completed")
public void markOrderCompleted(PaymentCompletedEvent event) {
    // update order status to COMPLETED
}
```

◆ Inventory Service

```
java

@KafkaListener(topics = "order-created")
public void reserveInventory(OrderCreatedEvent event) {
    // Check & reserve inventory
    kafkaTemplate.send("inventory-reserved", new InventoryReservedEvent(event.getOrderId()));
}
```

◆ Payment Service

```
java

@KafkaListener(topics = "inventory-reserved")
public void processPayment(InventoryReservedEvent event) {
    // Deduct money
    kafkaTemplate.send("payment-completed", new PaymentCompletedEvent(event.getOrderId()));
}
```

2. Orchestration-Based Saga

Flow:

- A central **Orchestrator service** controls the steps and invokes each service.
- 💡 All services must expose APIs for compensate actions (cancel, release, refund).

Saga Orchestrator Logic:

```
java Copy  
  
@Service  
public class SagaService {  
  
    @Autowired OrderClient orderClient;  
    @Autowired PaymentClient paymentClient;  
    @Autowired InventoryClient inventoryClient;  
  
    public void startOrderSaga(OrderRequest request) {  
        try {  
            orderClient.createOrder(request);  
            paymentClient.processPayment(request);  
            inventoryClient.deductStock(request);  
            System.out.println("Order Saga Completed Successfully");  
        } catch (Exception e) {  
            System.out.println("Saga failed: " + e.getMessage());  
            compensate(request);  
        }  
    }  
  
    private void compensate(OrderRequest request) {  
        try { inventoryClient.restoreStock(request); } catch (Exception ignored) {}  
        try { paymentClient.refundPayment(request); } catch (Exception ignored) {}  
        try { orderClient.cancelOrder(request); } catch (Exception ignored) {}  
    }  
}
```

Expose a REST Endpoint:

```
java Copy  
  
@RestController  
@RequestMapping("/saga")  
public class SagaController {  
  
    @Autowired SagaService sagaService;  
  
    @PostMapping("/order")  
    public ResponseEntity<String> createOrder(@RequestBody OrderRequest request) {  
        sagaService.startOrderSaga(request);  
        return ResponseEntity.ok("Saga Execution Started");  
    }  
}
```

2. 📦 Order Service

Endpoints:

```
java Copy  
  
@RestController  
@RequestMapping("/orders")  
public class OrderController {  
  
    @PostMapping  
    public ResponseEntity<String> createOrder(@RequestBody OrderRequest request) {  
        System.out.println("Order Created: " + request.getOrderId());  
        return ResponseEntity.ok("Order Created");  
    }  
  
    @PostMapping("/cancel")  
    public ResponseEntity<String> cancelOrder(@RequestBody OrderRequest request) {  
        System.out.println("Order Cancelled: " + request.getOrderId());  
        return ResponseEntity.ok("Order Cancelled");  
    }  
}
```

3. 💳 Payment Service

Endpoints:

```
java Copy  
  
@RestController  
@RequestMapping("/payments")  
public class PaymentController {  
  
    @PostMapping  
    public ResponseEntity<String> processPayment(@RequestBody OrderRequest request) {  
        System.out.println("Payment Processed for Order: " + request.getOrderId());  
        return ResponseEntity.ok("Payment Processed");  
    }  
  
    @PostMapping("/refund")  
    public ResponseEntity<String> refundPayment(@RequestBody OrderRequest request) {  
        System.out.println("Payment Refunded for Order: " + request.getOrderId());  
        return ResponseEntity.ok("Payment Refunded");  
    }  
}
```

4. 📦 Inventory Service

Endpoints:

```
java Copy

@RestController
@RequestMapping("/inventory")
public class InventoryController {

    @PostMapping("/deduct")
    public ResponseEntity<String> deductStock(@RequestBody OrderRequest request) {
        System.out.println("Stock Deducted: " + request.getProductId());
        // Simulate failure
        if ("fail".equals(request.getProductId())) {
            throw new RuntimeException("Stock not available");
        }
        return ResponseEntity.ok("Stock Deducted");
    }

    @PostMapping("/restore")
    public ResponseEntity<String> restoreStock(@RequestBody OrderRequest request) {
        System.out.println("Stock Restored: " + request.getProductId());
        return ResponseEntity.ok("Stock Restored");
    }
}
```

5. ⚪ Orchestrator Service

Feign Clients:

java

```
@FeignClient(name = "order-service")
public interface OrderClient {
    @PostMapping("/orders")
    void createOrder(@RequestBody OrderRequest request);

    @PostMapping("/orders/cancel")
    void cancelOrder(@RequestBody OrderRequest request);
}

@FeignClient(name = "payment-service")
public interface PaymentClient {
    @PostMapping("/payments")
    void processPayment(@RequestBody OrderRequest request);

    @PostMapping("/payments/refund")
    void refundPayment(@RequestBody OrderRequest request);
}
```

```
@FeignClient(name = "payment-service")
public interface PaymentClient {
    @PostMapping("/payments")
    void processPayment(@RequestBody OrderRequest request);

    @PostMapping("/payments/refund")
    void refundPayment(@RequestBody OrderRequest request);
}

@FeignClient(name = "inventory-service")
public interface InventoryClient {
    @PostMapping("/inventory/deduct")
    void deductStock(@RequestBody OrderRequest request);

    @PostMapping("/inventory/restore")
    void restoreStock(@RequestBody OrderRequest request);
}
```

43) Idempotency Situation ?

→ Idempotency means that performing the same operation multiple times produces the same result as doing it once.

📌 Examples:

Action	Idempotent?	Why?
GET /orders/123	<input checked="" type="checkbox"/> Yes	Always returns the same order data.
POST /orders	<input type="checkbox"/> No	Each call may create a new order (duplicate).
DELETE /orders/123	<input checked="" type="checkbox"/> Yes	Deleting the same order repeatedly has the same effect.

📌 Why is Idempotency Important?

In distributed systems (like microservices), network retries, timeouts, or client retries may cause the same request to be sent multiple times. Without idempotency:

- Duplicate entries might be created.
- Payments may be processed more than once.
- Data consistency may break.

→ Make non-idempotent operations (like POST) **idempotent** by handling duplicate requests gracefully.

How to Implement Idempotency in Spring Boot

Strategy

Use an **Idempotency Key** (a unique request identifier) sent in the header or body. Before processing the request:

1. Check if a request with that key has already been processed.
2. If yes, return the stored response.
3. If not, process it, store the key and result.

Components Needed

- A table to store processed idempotency keys.
- A filter or interceptor to check the key before processing.
- Custom logic to cache and return previous responses.

Step-by-Step Code

```
@Entity  
public class IdempotencyRecord {  
    @Id  
    private String idempotencyKey;  
  
    private String responseBody; // Optional: store serialized response  
    private LocalDateTime createdAt;  
}
```

```
public interface IdempotencyRecordRepository extends JpaRepository<IdempotencyRecord, String> {  
}
```

```
@RestController  
@RequestMapping("/orders")  
public class OrderController {  
  
    @Autowired  
    private OrderService orderService;  
  
    @PostMapping  
    public ResponseEntity<String> createOrder(  
        @RequestHeader("Idempotency-Key") String idempotencyKey,  
        @RequestBody OrderRequest orderRequest) {  
  
        return orderService.createOrder(orderRequest, idempotencyKey);  
    }  
}
```

```
@Service
public class OrderService {

    @Autowired
    private IdempotencyRecordRepository idempotencyRepo;

    public ResponseEntity<String> createOrder(OrderRequest request, String key) {
        Optional<IdempotencyRecord> recordOpt = idempotencyRepo.findById(key);

        if (recordOpt.isPresent()) {
            // Already processed, return stored response
            return ResponseEntity.ok("Duplicate request. Order already created.");
        }

        // ✅ Process business logic
        System.out.println("Processing order: " + request.getOrderId());

        // Save idempotency record
        IdempotencyRecord record = new IdempotencyRecord();
        record.setIdempotencyKey(key);
        record.setCreatedAt(LocalDateTime.now());
        idempotencyRepo.save(record);

        return ResponseEntity.ok("Order created successfully.");
    }
}
```

🔒 Bonus: Ensure Uniqueness at Database Level

To prevent race conditions:

sql

```
CREATE UNIQUE INDEX idx_idempotency_key ON idempotency_record(idempotency_key);
```

🧪 Testing the Idempotency

- Send the same request twice with the same `Idempotency-Key` header.
- First call processes the request.
- Second call returns the cached/stored response or a message.

💡 Best Practices

- Expire idempotency keys after a defined time (e.g., 24 hours).
- Only make `POST`, `PUT`, or `PATCH` idempotent.
- Use UUID or a hash of the request body as the idempotency key.

➔ The **idempotency key** is typically generated by the client (frontend/UI) and sent with each request to the backend.

Who Sends the Idempotency Key?

Option	Who Generates?	Notes
<input checked="" type="checkbox"/> Client (UI / Mobile / API consumer)	Generates and sends a unique key per action	Most common and scalable
<input type="checkbox"/> Backend	Accepts request without key and generates one internally	Rare — not useful for retries since UI doesn't know the key

Alternative: Backend Generates Key and Returns It

If your use case requires it:

1. Client sends a request without the key.
2. Backend generates a key, stores it, returns it.
3. Client uses that key for retries.

But this is less common and harder to manage in distributed systems.

Why Should the Client Send It?

1. To support retries safely (e.g., on timeouts or failures).
2. Backend can't detect retries if the request looks new each time.
3. Backend can trust and verify duplicate keys if client controls them.

Example: UI Sending Key

JavaScript / Axios Example

```
js

const idempotencyKey = crypto.randomUUID() // Or use UUID Library

axios.post('/orders', orderData, {
  headers: {
    'Idempotency-Key': idempotencyKey
  }
});
```

44) Discuss Spring Boot's approach to database migrations. How can one integrate tools like Flyway or Liquibase?

- Spring Boot integrates seamlessly with database migration tools like Flyway and Liquibase, automating the process of database schema creation and updates.
- To use Flyway, add its starter dependency to your pom.xml or build.gradle.
- Spring Boot automatically detects Flyway migrations in **src/main/resources/db/migration** and applies them accordingly.
- Spring Boot automatically detects Flyway migrations in **src/main/resources/db-changelog.sql** and applies them accordingly.

45) How do you implement caching in Spring Boot for improved performance?

💡 What is Caching?

Caching is storing data temporarily to serve future requests faster without hitting the database or a slow service every time.

🚀 Levels of Caching

Level	Example	Description
In-memory	@Cacheable , ConcurrentHashMap	Local to the app
Distributed	Redis, Hazelcast	Shared across microservices
Persistent	CDN, HTTP cache headers	For static assets

💻 Why Use Caching?

- Reduce database or API calls
- Improve response times
- Handle high traffic efficiently
- Save computational cost

Step-by-Step: Basic Caching in Spring Boot

◆ Step 1: Add Dependency

```
xml

<!-- Add in your pom.xml -->
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-cache</artifactId>
</dependency>
```

For Redis (later):

```
xml

<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-data-redis</artifactId>
</dependency>
```

◆ Step 2: Enable Caching

```
java

@SpringBootApplication
@EnableCaching
public class MyApp {
    public static void main(String[] args) {
        SpringApplication.run(MyApp.class, args);
    }
}
```

◆ Step 3: Use `@Cacheable`

```
java

@Service
public class ProductService {

    @Cacheable(value = "products", key = "#id")
    public Product getProductId(Long id) {
        simulateSlowService(); // e.g., DB or external call
        return productRepository.findById(id).orElse(null);
    }

    private void simulateSlowService() {
        try { Thread.sleep(3000); } catch (InterruptedException e) {}
    }
}
```

⊕ Explanation

- `@Cacheable`: Checks the cache before executing the method.
- `value`: Cache name.
- `key`: Key for the cache entry (`#id` = method parameter).

5 Other Cache Annotations

Annotation	Use-case
@Cacheable	Read from cache or method
@CachePut	Always run method & update cache
@CacheEvict	Remove from cache

```
@CacheEvict(value = "products", key = "#id")
public void deleteProduct(Long id) {
    productRepository.deleteById(id);
}
```

46) What are the best practices for managing application properties and secrets in Spring Boot?

→ Managing application properties and secrets in Spring Boot is crucial for security, maintainability, and environment-specific configuration.

✳ 1. Organize application.properties or application.yml

Split configurations by environment:

```
arduino

src/
└── main/
    └── resources/
        ├── application.yml      # default
        ├── application-dev.yml  # dev-specific
        └── application-prod.yml # prod-specific
```

🔒 2. Never Hardcode Secrets

Bad practice:

```
yaml

db:
    username: admin
    password: mySecret123 ✗
```

✓ 3. Use Environment Variables

Let external systems pass secrets:

```
yaml

spring:
  datasource:
    username: ${DB_USER}
    password: ${DB_PASS}
```

Then set in your environment:

```
bash

export DB_USER=admin
export DB_PASS=superSecret
```

⌚ 4. Use @ConfigurationProperties for Structured Config

Encapsulate grouped settings:

```
java

@ConfigurationProperties(prefix = "payment")
@Component
public class PaymentProperties {
    private String apiKey;
    private String merchantId;
}
```

```
yaml

payment:
  apiKey: ${PAYMENT_API_KEY}
  merchantId: 12345
```

5. Use a Secret Manager for Production

Use services like:

- AWS Secrets Manager
- Azure Key Vault
- HashiCorp Vault
- Spring Cloud Config with encryption

Example: Spring Cloud Vault

```
yaml

spring:
  cloud:
    vault:
      uri: http://localhost:8200
      token: ${VAULT_TOKEN}
```

6. Secure application.yml Files

- Never commit secrets to Git
- Add `application-prod.yml` to `.gitignore` if it contains sensitive data
- Use `.env` files only for local dev, not for production

7. Use Profiles with @Profile Annotation

Activate beans per environment:

```
java

@Profile("dev")
@Bean
public DataSource devDataSource() { ... }

@Profile("prod")
@Bean
public DataSource prodDataSource() { ... }
```

✓ Summary: Best Practices

Practice	Benefit
Use profiles (<code>application-dev.yml</code>)	Clean separation per environment
Store secrets in env vars or Vault	Secure, scalable, CI/CD friendly
Use <code>@ConfigurationProperties</code>	Organized, type-safe configs
Validate properties	Avoid runtime config errors
Use encryption for sensitive data	Security for in-code secrets
Avoid hardcoding or committing secrets	Prevent leaks

47) What are the tools available for detecting memory leak ?

VisualVM

- VisualVM is an all-in-one Java troubleshooting tool that integrates several JDK command-line tools .
- It is lightweight performance and memory profiling capabilities.
- It's included in the Oracle JDK download.

Key Features:

- Monitor application memory consumption in real-time.
- Analyze heap dumps to identify memory leaks.
- Track down memory leaks with its built-in heap walker.

Eclipse Memory Analyzer (MAT)

- Eclipse MAT is a specialized tool designed for analyzing heap dumps.
- It is particularly effective in identifying memory leaks and reducing memory consumption.

Key Features:

- Analyze large heap dumps.
- Automatically identify memory leak suspects.
- Provide detailed reports on memory consumption by objects.

Usage Example:

**After obtaining a heap dump from a running application
(which can be triggered in the JVM on OutOfMemoryError),
MAT can be used to analyze this dump.**

**It provides a histogram of objects in memory, allowing developers to see which classes and objects are
consuming the most memory.**

JProfiler

- JProfiler is a comprehensive profiling tool for Java with capabilities for both memory and performance profiling.
- It's a commercial tool but is widely regarded for its user-friendly interface and detailed analysis.

Key Features:

- Real-time memory and CPU profiling.
- Advanced heap analysis and visualization.
- Ability to track every object in the heap and analyze memory consumption.

48) How do you implement a custom health check in Spring Boot Actuator?

- To implement a custom health check in Spring Boot Actuator, you need to create a bean that implements the **HealthIndicator** interface.
- Spring Boot will automatically include your custom health check in the /actuator/health endpoint.

```
@Component
public class MyCustomHealthIndicator implements HealthIndicator {

    @Override
    public Health health() {
        // Perform some custom logic (e.g., DB check, external service ping)
        boolean isServiceUp = checkCustomServiceHealth();

        if (isServiceUp) {
            return Health.up()
                .withDetail("customService", "Available")
                .build();
        } else {
            return Health.down()
                .withDetail("customService", "Not reachable")
                .build();
        }
    }

    private boolean checkCustomServiceHealth() {
        // Replace with real check
        return true;
    }
}
```

3. Enable Health Details in application.yml

```
yaml

management:
  endpoints:
    web:
      exposure:
        include: health,info
  endpoint:
    health:
      show-details: always
```

49) How would you handle a situation where a Spring Boot application runs out of memory?

→ When a Spring Boot application runs out of memory (OOM), it's often due to memory leaks, inefficient data structures, or misconfigured JVM settings.

🛠 Step 1: Identify the Root Cause

🔍 Check the error logs:

Look for `java.lang.OutOfMemoryError` in logs:

- `OutOfMemoryError: Java heap space`
- `OutOfMemoryError: GC overhead limit exceeded`
- `OutOfMemoryError: Metaspace`

📈 Enable heap dump on OOM:

Add to JVM options:

```
bash
-XX:+HeapDumpOnOutOfMemoryError -XX:HeapDumpPath=/path/to/dumps
```

Analyze heap dump with:

- [Eclipse MAT \(Memory Analyzer Tool\)](#)
- VisualVM
- JProfiler

📝 Step 2: Analyze Memory Usage

Use runtime tools:

- `jstat`, `jmap`, `jstack`, `jconsole`
- Spring Boot Actuator: enable `/actuator/heapdump`, `/metrics`

⚙️ Step 3: Configure Appropriate JVM Settings

Set appropriate heap size in `application startup` or Docker config:

```
bash
-Xms512m -Xmx2g
```

Also consider:

```
bash
-XX:MaxMetaspaceSize=256m
```

Step 4: **Code-Level Fixes**

1. Avoid loading large datasets into memory:

```
java

// ✗ Avoid
List<Product> products = productRepository.findAll();

// ✓ Use pagination or streaming
Page<Product> products = productRepository.findAll(PageRequest.of(0, 100));
```

50) How would you design a Spring Boot application to handle millions of concurrent users?

→ Designing a Spring Boot application to handle millions of concurrent users requires careful planning across architecture, scalability, performance optimization, and fault tolerance.

1. Use a Microservices or Modular Monolith Architecture

Break the application into functional services:

- User Service
- Product/Catalog Service
- Order Service
- Payment Service
- Inventory Service
- Notification Service

Benefits:

- Scalability per feature
- Fault isolation
- Easier deployments

3. Scalability Strategy

Horizontal Scaling:

- Use containers (Docker) and orchestration (Kubernetes)
- Autoscale with Horizontal Pod Autoscaler (HPA) based on CPU/memory/requests

Stateless Services:

- Design services to be **stateless** so they can scale easily
- Store session state in **Redis** or another distributed cache

2. API Gateway and Load Balancer

- Use Spring Cloud Gateway or Kong/Nginx as the API Gateway
- Use AWS ELB, Azure Application Gateway, or GCP Load Balancing for traffic distribution

Key Features:

- Rate limiting
- IP filtering
- Authentication/authorization
- Routing and versioning

4. Database Design for High Throughput

- Use sharded or partitioned databases
- Read/write separation (replication)
- Connection pooling (HikariCP)
- Use NoSQL (e.g. Cassandra, DynamoDB) for high-velocity data

Optimize:

- Indexes
- Query execution plans
- Avoid joins and heavy aggregations in high-throughput paths

5. Use Caching Aggressively

Reduce DB hits with:

- `@Cacheable` + Caffeine/Redis
- CDN for static content (CloudFront, Akamai)

Example:

```
java  
  
@Cacheable(value = "products", key = "#id")  
public Product getProductById(Long id) { ... }
```

6. Asynchronous Processing

Offload non-critical workloads using:

- Kafka, RabbitMQ, or AWS SQS
- Use Spring Boot with `@Async`, `@Scheduled` for background tasks

Example use cases:

- Email notifications
- Order processing
- Event logging

7. Security and Rate Limiting

- Use JWT + OAuth2 for authentication
- Use API Gateway or filters for throttling
- Protect endpoints with Spring Security and custom filters



9. Resilience and Failover

- Use **Resilience4j** for:
 - Circuit breakers
 - Rate limiters
 - Retry logic
- Use **K8s readiness/liveness probes**
- Geo-redundant deployment (multi-AZ, multi-region)
- Graceful fallback mechanisms



10. Performance Tuning

- Tune JVM parameters (`-Xmx`, GC, etc.)
- Use **non-blocking WebClient** (instead of RestTemplate)
- Tune thread pools and database connection pools
- Compress responses (`gzip`)
- Batch expensive operations

51) What are the best practices for logging in Spring Boot?

1. Use a Standard Logging Framework

- Spring Boot uses **SLF4J** as the logging facade by default.
- Common implementations:
 - **Logback** (default)
 - **Log4j2** (optional)
- Stick to **SLF4J API** in your code for flexibility.

java

```
private static final Logger logger = LoggerFactory.getLogger(MyClass.class)
logger.info("This is a log message");
```

2. Configure Logging Levels Appropriately

- Use levels **TRACE, DEBUG, INFO, WARN, ERROR** judiciously.
- Avoid DEBUG/TRACE in production unless troubleshooting.
- Configure levels in `application.properties` OR `application.yml`.

properties

```
logging.level.org.springframework.web=DEBUG
logging.level.com.myapp=INFO
```

3. Externalize Log Configuration

- Use `logback-spring.xml` or `log4j2-spring.xml` for advanced config.
- Avoid hardcoding settings in code.
- Enable easy changes without redeploying.

8. Log Exceptions Properly

- Always log the exception stack trace, not just the message.

java

```
try {
    // risky code
} catch (Exception e) {
    logger.error("Error processing request", e);
}
```

52) What are the best practices for securing a Spring Boot application?

→ Securing a Spring Boot application involves multiple layers—from authentication and authorization to protecting against common vulnerabilities and securing infrastructure.

1. Use Spring Security

- Leverage **Spring Security**, the de-facto security framework for Spring apps.
- Configure authentication and authorization with roles and permissions.
- Use **OAuth2 / OpenID Connect** with providers like Keycloak, Okta, or Auth0 for SSO.
- Prefer **stateless JWT tokens** for REST APIs.

2. Secure Authentication

- Store passwords securely using **bcrypt** or other strong hashing algorithms.
- Implement **multi-factor authentication (MFA)** if needed.
- Limit login attempts to prevent brute force attacks.

3. Secure Authorization

- Use **method-level security** (`@PreAuthorize`, `@Secured`) to restrict sensitive operations.
- Enforce **principle of least privilege**: users get only the permissions they need.
- Validate permissions both on frontend and backend.

5. Use HTTPS

- Always serve your app over **HTTPS** with valid certificates.
- Redirect HTTP traffic to HTTPS.

4. Protect Against Common Vulnerabilities

- **SQL Injection**: Use **parameterized queries** or Spring Data repositories instead of string concatenation.
- **Cross-Site Scripting (XSS)**: Sanitize inputs, use frameworks' escaping features, enable Content Security Policy (CSP).
- **Cross-Site Request Forgery (CSRF)**: Enable CSRF protection on web forms (Spring Security enables it by default).
- **Clickjacking**: Use HTTP headers like `X-Frame-Options` or Spring Security's defaults.
- **Secure Headers**: Set HTTP security headers (`Strict-Transport-Security`, `X-Content-Type-Options`, etc.)

53) How do you ensure code quality in a Spring Boot project?

1. Follow Clean Code Principles

- Use meaningful variable and method names.
- Keep methods small and focused on a single responsibility.
- Avoid duplicate code (DRY principle).
- Structure code into logical layers: Controller, Service, Repository.

2. Use Static Code Analysis Tools

- Integrate tools like:
 - SonarQube – for code smells, bugs, and coverage.
 - Checkstyle – for enforcing coding standards.
 - PMD/FindBugs (SpotBugs) – for detecting bugs and anti-patterns.
- Run analysis in CI pipelines for automatic feedback.

3. Write Unit Tests

- Use **JUnit 5** and **Mockito** to test services and utility classes.
- Aim for high coverage, but focus on meaningful tests.

5. Use Code Reviews

- Enforce pull requests (PRs) with mandatory code reviews.
- Review code for:
 - Logic correctness
 - Code clarity
 - Performance
 - Security
 - Test coverage

4. Write Integration Tests

- Use **Spring Boot Test** (`@SpringBootTest`, `@WebMvcTest`) to test APIs and database interactions.
- Run tests against in-memory databases (e.g., H2) for speed.

6. Adopt CI/CD with Quality Gates

- Use CI tools (e.g., GitHub Actions, Jenkins, GitLab CI) to:
 - Run tests
 - Lint code
 - Run static analysis
- Use SonarQube quality gates to block merging code with issues.

7. Use Consistent Code Formatting

- Use IDE formatters or tools like **Spotless** or **Prettier** for Java.
- Enforce formatting using pre-commit hooks or CI pipelines.

8. Document the Code

- Use Javadoc for public APIs.
- Document assumptions, edge cases, and important business logic.

10. Monitor Code Metrics

Track:

- Code coverage
- Code duplication
- Cyclomatic complexity
- Number of issues over time

55) How do you implement OAuth2 authentication in Spring Boot?

Steps to Implement OAuth2 Client Authentication in Spring Boot

OAuth2 Scenarios

You can configure Spring Boot to:

1. Authenticate users via an **OAuth2 provider** (Google, GitHub, Keycloak).
2. Secure your REST APIs using **JWT tokens** (resource server).
3. Act as an **authorization server** (with Spring Authorization Server).

2. Configure `application.yml` or `application.properties`

Example: Login with Google

```
yaml

spring:
  security:
    oauth2:
      client:
        registration:
          google:
            client-id: YOUR_GOOGLE_CLIENT_ID
            client-secret: YOUR_GOOGLE_CLIENT_SECRET
            scope: profile, email
        provider:
          google:
            authorization-uri: https://accounts.google.com/o/oauth2/v2/auth
            token-uri: https://oauth2.googleapis.com/token
            user-info-uri: https://www.googleapis.com/oauth2/v3 userinfo
```

1. Add Dependencies

```
xml

<!-- OAuth2 Client -->
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-oauth2-client</artifactId>
</dependency>
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-security</artifactId>
</dependency>
```

3. Enable OAuth2 Login

No extra configuration is required. Spring Boot auto-configures a login page at `/login`.

4. Protect Routes (Optional Custom Config)

```
java

@Configuration
@EnableWebSecurity
public class SecurityConfig {

    @Bean
    public SecurityFilterChain securityFilterChain(HttpSecurity http) throws Exception {
        http
            .authorizeHttpRequests(auth -> auth
                .requestMatchers("/", "/public").permitAll()
                .anyRequest().authenticated()
            )
            .oauth2Login() // Enable OAuth2 login
        return http.build();
    }
}
```

Copy

5. Access Authenticated User Details

```
java

@GetMapping("/profile")
public String userDetails(@AuthenticationPrincipal OAuth2User principal) {
    return "Hello " + principal.getAttribute("name");
}
```

Securing REST APIs (OAuth2 Resource Server)

1. Add Dependency

```
xml  
  
<dependency>  
  <groupId>org.springframework.boot</groupId>  
  <artifactId>spring-boot-starter-oauth2-resource-server</artifactId>  
</dependency>
```

2. Configure JWT Issuer (e.g., Keycloak, Auth0, Okta)

```
yaml  
  
spring:  
  security:  
    oauth2:  
      resourceserver:  
        jwt:  
          issuer-uri: https://your-auth-domain/.well-known/openid-configuration
```

3. Secure Endpoints

```
java  
  
@Configuration  
@EnableWebSecurity  
public class SecurityConfig {  
  
  @Bean  
  public SecurityFilterChain securityFilterChain(HttpSecurity http) throws Exception {  
    http  
      .authorizeHttpRequests(auth -> auth  
        .anyRequest().authenticated()  
      )  
      .oauth2ResourceServer()  
      .jwt();  
    return http.build();  
  }  
}
```

Optional: Role-Based Access Control

Use JWT claim mappings or Keycloak roles:

```
java  
  
@PreAuthorize("hasRole('ADMIN')")  
@GetMapping("/admin")  
public String adminOnly() {  
  return "Admin content";  
}
```

56) What is the difference between Basic Authentication and JWT in Spring Boot?

→ The difference between Basic Authentication and JWT (JSON Web Token) in Spring Boot lies in how authentication information is transmitted, stored, and managed.

1. Basic Authentication

✓ How It Works

- Client sends username and password in the `Authorization` header of every HTTP request:

pgsql

Copy

```
Authorization: Basic base64(username:password)
```

- Server validates credentials on **every request**.

⚠ Risks

- Less secure, especially over unsecured (non-HTTPS) connections.
- Can't easily revoke access without changing password.

✓ Characteristics

Feature	Description
State	Stateless (but credentials sent every time)
Credential Reuse	Password is sent with every request (security risk)
Token Format	No token; base64-encoded credentials
Expiration	None (unless managed externally)
Server Storage	Not needed
Use Case	Simple internal services, basic auth testing

2. JWT (JSON Web Token) Authentication

✓ How It Works

- User authenticates once → server issues a **JWT** (digitally signed token).
- Client sends the token in the `Authorization` header:

```
makefile  
  
Authorization: Bearer eyJhbGciOiJIUzI1NiIsInR5cCI...
```
- Server verifies the token's signature and claims on each request.

🔍 Feature Comparison

Feature	Basic Authentication	JWT Authentication
Auth Header	<code>Authorization: Basic</code>	<code>Authorization: Bearer</code>
Credential Sent	Every request (username/password)	Once (token issued)
Stateless	Yes	Yes
Security Risk	High (password reuse)	Lower (no password in each request)
Revocation	Change password	Harder; needs token blacklist/cache
Use with OAuth2	Not compatible	Standard approach
Token Content	None	Claims like user ID, roles, expiry
Best Use	Dev, test, simple internal apps	Production APIs, scalable services

✓ Characteristics

Feature	Description
State	Fully stateless
Credential Reuse	Token is reused (not password)
Token Format	Self-contained (contains user info, roles, expiry)
Expiration	Yes (usually short-lived)
Server Storage	No need (unless using blacklists)
Use Case	Scalable APIs, mobile/backend auth, microservices

🛠 Spring Boot Implementation

- **Basic Auth:** Use `HttpBasic` in security config

```
java  
  
http  
    .authorizeHttpRequests().anyRequest().authenticated()  
    .and().httpBasic();
```

- **JWT:** Use `spring-boot-starter-oauth2-resource-server`

```
yaml  
  
spring:  
  security:  
    oauth2:  
      resourceserver:  
        jwt:  
          issuer-uri: https://example.com/auth ↴
```

57) How do you implement role-based access control (RBAC) in Spring Boot?

58) How do you optimize database queries in Spring Boot using JPA or Hibernate?

1. Use Projections Instead of Fetching Entire Entities

Fetch only the required fields using DTO projections.

Bad:

```
java
List<User> users = userRepository.findAll(); // Fetches entire entity
```

Optimized:

```
java
public interface UserNameEmailProjection {
    String getName();
    String getEmail();
}
List<UserNameEmailProjection> findByActiveTrue();
```

2. Use Pagination for Large Result Sets

Avoid returning thousands of rows at once.

```
java
Page<User> findByStatus(String status, Pageable pageable);
```

```
java
Pageable pageable = PageRequest.of(0, 20);
Page<User> users = userRepository.findByStatus("ACTIVE", pageable);
```

3. Avoid N+1 Problem Using Fetch Joins

N+1 happens when you fetch a list of entities and each entity triggers additional queries for related entities.

 Bad:

```
java Copy Edit
List<Order> orders = orderRepository.findAll(); // Then multiple SELECTs for each order items
```

 Optimized with Join Fetch:

```
java Copy Edit
@Query("SELECT o FROM Order o JOIN FETCH o.items")
List<Order> findAllWithItems();
```

Or annotate entity with:

```
java
@OneToOne(fetch = FetchType.LAZY)
@JoinFetch
private List<Item> items;
```

8. Use Native Queries for Complex Joins

If JPA queries are too slow or unreadable:

```
java Copy
@Query(value = "SELECT * FROM users WHERE status = 'ACTIVE'", nativeQuery = true)
List<User> findActiveUsers();
```

7. Batch Inserts/Updates

Reduce number of insert/update queries:

```
yaml
spring.jpa.properties.hibernate.jdbc.batch_size=30
spring.jpa.properties.hibernate.order_inserts=true
spring.jpa.properties.hibernate.order_updates=true
```

entity:

```
java
@BatchSize(size = 30)
```

10. Cache Frequently Accessed Data

Use 2nd level cache or Spring Cache abstraction:

```
java
```

```
@Cacheable("users")
public User getUserById(Long id) { ... }
```

Summary of Optimization Techniques

Technique	Benefit
DTO/Projections	Reduces memory and data transfer
Pagination	Handles large datasets efficiently
Fetch Joins & <code>@EntityGraph</code>	Avoids N+1 queries
SQL Logging	Helps identify slow or inefficient SQL
Batch Operations	Reduces DB round trips
Indexes	Speeds up queries
Native Queries	More control and performance
Caching	Reduces DB load for frequently accessed data

58) How would you handle a situation where a Spring Boot application fails to start due to a circular dependency?

How to Handle Circular Dependency in Spring Boot

1. Refactor Your Code to Remove Circular Dependencies

- Best practice: redesign your classes so dependencies flow in one direction.
- Use interfaces and abstractions to break tight coupling.
- For example, introduce a third bean that both depend on instead of depending on each other.

3. Use Setter Injection Instead of Constructor Injection

Setter injection allows Spring to create beans first and then inject dependencies.

```
java

@Component
public class BeanA {
    private BeanB beanB;

    @Autowired
    public void setBeanB(BeanB beanB) {
        this.beanB = beanB;
    }
}

@Component
public class BeanB {
    private BeanA beanA;

    @Autowired
    public void setBeanA(BeanA beanA) {
```

2. Use `@Lazy` Annotation

Mark one or both dependencies as lazy to delay injection until needed, allowing Spring to create pr break the cycle.

```
java

@Component
public class BeanA {
    private final BeanB beanB;

    public BeanA(@Lazy BeanB beanB) {
        this.beanB = beanB;
    }
}

@Component
public class BeanB {
    private final BeanA beanA;

    public BeanB(@Lazy BeanA beanA) {
        this.beanA = beanA;
```

